

Mortar: 30-Minute Presentation Script

Prepared for the NexTechnologies-MY/Mortar project, 28 September 2026. You are the main presenter. This is one universal script: anyone joining the presentation should learn the complete argument and workflow.

Timing: 30 minutes of presentation. Questions follow afterward. **Use:** rehearse the wording, then speak from the short cue sheet. Do not read the slides. Time includes operating the prototype, pauses and transitions.

Presentation pages: use `presentation.html` or `presentation-slides.pdf` from this folder. The page numbers below refer to that new companion deck, not the older 20-page qualification deck in the repository. The mapping to the older deck is at the end of this document.

Working prototype: [Mortar on Render](#). All demo buyers, projects and banks are synthetic. Aster Heights is a demo project, not a verified Chin Hin development.

Agenda And Timing

Time	Page	Topic	Purpose
00:00–01:00	1	Business opening	Establish the outcome
01:00–01:30	2	Agenda	Explain the route through the presentation
01:30–04:00	3	Chin Hin company context	Connect the solution to this business
04:00–06:00	4	Evidence and unknowns	Explain what we know and how we know
06:00–07:30	5	Before: current workflow	Show the operational gap
07:30–09:30	6	Our solution logic	Explain how Mortar causes action
09:30–10:15	7 → prototype <code>/settings</code>	Demo boundaries	Identify synthetic data and current state
10:15–12:15	8 → prototype <code>/chase</code>	Daily work and task ownership	Demonstrate the first operational action
12:15–16:15	9 → prototype <code>/bookings/BK-9001</code>	Critical working workflow	Confirm a blocker, record receipt, show the change
16:15–18:00	10	Before and after	Explain exactly what improved
18:00–19:30	11 → prototype <code>/legal</code>	Legal handoff	Demonstrate the next department's work

19:30–21:00	12 → prototype <code>/manager</code>	Management intervention	Show accountable escalation
21:00–23:00	13 → prototype <code>/forecast</code>	Forecast and uncertainty	Explain what Finance can use
23:00–24:00	14	Technology second	Explain reliability after demonstrating value
24:00–26:00	15	Impact and measurement	Derive an example without claiming results
26:00–28:00	16 → prototype <code>/import</code>	Adoption	Show how daily use starts and continues
28:00–29:15	17	Built, proposed and unproven	Make the prototype's maturity clear
29:15–30:00	18	Close and request	Ask for a concrete pilot decision

01 — Business Opening

Time: 00:00–01:00. **Screen:** page 1, “A booking is a promise. A verified SPA is a conversion.” **Cue:** unit held → action → verified outcome.

Say:

“Imagine a project launches and the booking numbers look strong. Behind one of those bookings, the bank is still waiting for the buyer's latest payslips. Sales knows the buyer is interested. Loan Administration knows the file is incomplete. Legal is waiting for the case to reach them.

The business question is very practical: who takes the next action, and how do we know the booking actually moved forward?

We built Mortar to help the developer follow a booking through to a verified Sale and Purchase Agreement, or identify a case that needs an authorised release decision. Today I will show the working sequence: a blocker, an owner, a task, evidence of what happened, and the resulting change to the case.”

Delivery: look at the audience. Pause after “who takes the next action?” Do not open with the software stack or a list of AI features.

02 — Agenda

Time: 01:00–01:30. **Screen:** page 2, “Business → workflow → proof → adoption.”

Say:

“I will start with Chin Hin's business context and the evidence behind our choices. Then we will follow the workflow before and after Mortar, including the critical steps in the prototype. After that, I will explain the forecast, the

expected impact calculation, and how a twelve-week pilot could be adopted.

We have reserved the full thirty minutes for the presentation. Questions can follow afterward, and we can return to any case or calculation.”

03 — Why This Matters To Chin Hin

Time: 01:30–04:00. **Screen:** page 3. **Cue:** correct entity → buyer journey → existing systems → business implication.

Say:

“We researched Chin Hin at two levels. Chin Hin Group Berhad is the broader listed group. Chin Hin Group Property Berhad, or CHGP, is the property business most directly relevant to this challenge. Figures from those two reporting entities must be labelled separately.

CHGP's 2025 annual report describes a focus on first-time owners and urban residential development. It reports approximately RM976.7 million in annual revenue and RM2.18 billion in unbilled sales at year-end. Those are reported company figures. They establish the scale and commercial context; they do not tell us how many unsigned bookings are stuck.

Unbilled sales are contracted sales awaiting future recognition as projects progress. They are not the same as the unsigned booking pool Mortar addresses. Likewise, an SPA signing is not receipt of the entire property price.

The company's own first-homebuyer guide puts checking loan options and preparing documents before signing agreements. That supports focusing on the handoff between buyer documentation, bank assessment and legal execution.

There is also an important system context. Chin Hin announced a Kingdee ERP transformation in July 2026. So our integration conversation must start with its actual operational systems and approved exports. IFCA appears in one practitioner's survey response; it is not evidence that Chin Hin uses IFCA.

Our interpretation is that Mortar should prove a small, useful follow-up workflow within that existing environment. The sponsor's brief describes fragmented booking information, but we have not audited Chin Hin's live files. We would validate the gap with its staff before deciding where Mortar belongs.

We therefore propose starting with one project and one daily administrator, then testing whether confirmed follow-up produces more timely, verified SPAs.”

Evidence: [CHGP Annual Report 2025](#), printed pp. 2–3; [first-homebuyer guide](#); [One Chin Hin transformation, 29 July 2026](#). Read `company-research.md` for the fuller briefing and source distinctions.

04 — Evidence, Assumptions And Unknowns

Time: 04:00–06:00. **Screen:** page 4. **Cue:** brief ≠ survey ≠ simulation.

Say:

"We use four evidence categories. First, the sponsor's challenge brief describes the problem it wants addressed. Second, official company publications provide business context. Third, our industry practitioner survey helps us prioritise what to investigate. Fourth, the prototype uses simulated cases to demonstrate the mechanics.

The current survey has eight anonymous respondents, collected between the eighteenth and twenty-second of September. Seven named loan rejection or insufficient financing as the single biggest leakage cause. All eight placed it in their top three. Five named bank credit assessment as the stage where cases get stuck. Seven wanted an earlier financing eligibility check.

This is a small convenience sample with mixed roles. Only one respondent covered loan administration and legal work. These are counts of respondents, not percentages of failed Chin Hin bookings. We cannot conclude that seven out of eight Chin Hin failures are caused by financing.

The practitioner with loan and legal experience prioritised slow bank processing. That matters: we must test both financing problems and delayed follow-up, rather than assume one explanation covers every project.

Our own logic is to make financing concerns visible and chase recoverable blockers early. We cannot overturn a bank's credit decision. We can help staff complete documents, follow up, review alternative applications and escalate a genuinely stalled case.

The missing evidence is Chin Hin's actual leakage rate, cause distribution, staff workload and ordinary follow-up performance. Our pilot is designed to collect that evidence. The blank interview template is not a completed interview and we do not cite it as one."

Source: repository `docs/research/practitioner-survey/README.md` and `docs/source/problem-statement.md`. Say "7 of 8 respondents," never "87.5% of Chin Hin's bookings."

05 — Before: The Workflow Gap

Time: 06:00–07:30. **Screen:** page 5. **Cue:** update arrives → responsibility unclear → no verified next move.

Say:

"This is the workflow described by the brief, expressed as a concrete case. A booking exists. The buyer and banker exchange messages. A document request arrives. Someone has to identify the requirement, decide who should obtain it, and make sure the reply gets back into the booking record.

When those steps are split across messages and departmental notes, a file can remain marked 'with bank' even though the real next move belongs to the buyer or an internal administrator. The important detail is who can act now.

The delay continues if no one owns the chase, no due date is set, or a reply is not recorded. Legal may receive a case late, and management may see a booking without understanding how much of the conversion journey remains.

This diagram is a model of the challenge, not a measured time-and-motion study of Chin Hin. We would check it with Sales, Loan Administration and Legal. The change we propose is to make the next move and its evidence explicit."

Action: trace the diagram with the pointer. Do not introduce invented current-process timings or say Chin Hin has no existing system.

06 — Our Solution Logic

Time: 07:30–09:30. **Screen:** page 6. **Cue:** evidence → blocker → owner → task → confirmation → updated case.

Say:

“Mortar starts with a case record for each booking. It tracks loan and legal events separately, because those activities can overlap. It also keeps bank applications separate, so one rejected application does not automatically mean the whole booking has failed.

Our first rule is that confirmed events drive the case state. A proposed interpretation remains visible for review, but it does not move the milestone. We retain when an event happened, when it was recorded, who reported it and who verified it.

Our second rule is to translate the case into the next practical action. If a document is outstanding, obtain it. If the complete application is waiting on a bank beyond the configured threshold, chase the banker. If financing is approved but no appointment is recorded, arrange the signing. If every reasonable recovery route has ended, bring the case to an authorised person for review.

Our third rule is to turn that action into work with an owner and a due date. Then the next reply must be recorded and verified. That closes the loop.

AI assists with interpreting short messages and suggesting actions. The business workflow remains explainable through the event record and the rules. An administrator can enter an update manually if the AI is unavailable or wrong.

This is our solution logic: finding a case is useful only if somebody takes the next step and the system can show what happened afterward. We will now demonstrate that loop on one synthetic booking.”

07 — Prepare The Audience For The Demo

Time: 09:30–10:15. **Screen:** page 7, then [Settings](#) as **Project Manager**.

Action: point to “Simulated Data,” seed and reference date. Do not press Add Demo Data or Delete Demo Data during the presentation. Switch through the top-right “Switch persona” menu to **Nurul Aina · Sales Admin** afterward.

Say:

“The interface is working software, and the cases are synthetic. This dataset uses a fixed reference date, the eighteenth of September, so today on the screen is the demo's day, not the calendar date of this presentation.

The numbers illustrate the workflow and calculation. They do not represent Chin Hin's records, and they do not prove commercial improvement. The named profiles demonstrate staff responsibilities and scoped case access;

they are not production identity verification.”

08 — Today: A Named Next Action

Time: 10:15–12:15. **Screen:** page 8 briefly, then [Today](#) as **Nurul Aina · Sales Admin**. **Case:** BK-9001, unit A-12-03, Raymond Tan Wei Hong.

Actions:

- Locate BK-9001 by its reference, not its position in the queue.
- Point to the bank-delay blocker and the default “Call The Banker” action.
- Point to Jev's alternative: “Ask For The Missing Document Instead.”
- If the payslip task already exists, point to “Task Open,” owner and due date. If it does not, select **Do That Instead** for the document chase and show the created task. Use **Create Task** only if its visible action matches the task you intend to demonstrate.
- Click the unit/buyer to open Quick View. Use its full-case link, or navigate directly to `/bookings/BK-9001`.

Say:

“This is the administrator's daily work. Each case tells us the blocker and the next step. On Raymond's case, the confirmed record says the bank has not decided. The default next step is to call the banker.

But there is also a banker message that suggests a more specific issue: missing payslips. Jev offers that as an alternative. We have deliberately kept the confirmed state and the proposed interpretation distinct.

Here is the action the system causes: an assigned document chase with a due date. [If present: This task is already open for Nurul Aina.] [If absent: I select the document action, and it becomes a task for the responsible staff member.]

Mortar records the work; the administrator still contacts the buyer through the existing approved channel. This button is not sending a WhatsApp message to the buyer.

Now let us inspect the evidence. We can open the case from this queue, so the administrator does not need to hunt through a separate departmental record.”

Hold: allow the audience to read the owner and due date for three seconds. Current live state on 28 September already contained a BK-9001 payslip task.

09 — Critical Workflow: A Message Becomes Confirmed Evidence

Time: 12:15–16:15. **Screen:** page 9 briefly, then [BK-9001](#). Remain **Nurul Aina · Sales Admin** for the main demonstration.

Actions and spoken cues:

- **Case header and tracks.** Point to "With Bank," the loan track and the unconfirmed Documents Requested event.

"The bank submission is confirmed. This document request is still a proposal. Notice that a confident AI reading has not become a fact."

- **Banker message.** Scroll to Kelvin Teo's message containing "latest 3 months slip gaji." Read the short relevant phrase, not the whole thread.

"The banker is asking for the latest three months' payslips. Jev proposes Documents Requested, Payslip, Sales. A person reviews that interpretation."

- **Confirm request.** Select **Confirm** on that message's proposal. Wait for the refreshed case. Point to Payslip Outstanding / the buyer-document wait and the now-confirmed event.

"I confirm the request. The case now reflects the missing document. The next action follows the evidence we just accepted."

Pause for three seconds. This is the first visible proof moment.

- **Simulated buyer reply.** Select **Paste A Message**. Leave Sender Role as **Buyer** and the synthetic buyer name. Paste:

Hi Nurul, slip gaji 3 bulan dah saya hantar ke Mei Ling tadi. Dokumen semua lengkap ya.

Leave the timestamp on the demo's reference day and select **Add Message**. Wait for the proposal; show the source tag, especially if it says cached or unavailable.

"For this demonstration, we simulate the buyer saying the payslips have reached Loan Administration. The system reads the reply, but a message claiming receipt still needs verification."

- **Verify receipt.** When Documents Received / Payslip appears, explain that the administrator has checked the receipt with Loan Administration in the demo scenario. Select **Confirm**. If the proposal is incorrect, dispute or dismiss it and use **Record An Update** → **Documents Received** → **Payslip** after verifying the simulated receipt. Choose **Apex Bank** when specifying the bank's document receipt; save with **Record Update**.

"In the operating workflow, staff check that the document actually arrived. We represent that check here, then confirm the receipt. An AI interpretation alone is insufficient."

- **Show the changed case.** Return to the header and loan track. The payslip should no longer be outstanding. The case still awaits the bank decision. Expand **Show Full History** to show the confirming actor and dates.

"The document blocker has cleared. We have not approved the mortgage. The bank still decides. What changed is the evidence, the next responsibility and the work required from the administrator."

Pause for three seconds. This is the second proof moment.

- **Complete work deliberately.** Show the document task. Select **Complete** only when the demonstrated document-chase obligation has been fulfilled. Explain that confirming a receipt and completing a task are separate acts.

"The task is completed after the work is done. The evidence remains in the case, so the next administrator can see why the blocker cleared."

Between actions, explain:

"We have shown a complete local loop: identify the issue, assign the chase, review the request, receive the reply, verify the receipt and update the case. No bank or solicitor has been required to adopt a new portal.

Notice also that low financing risk did not mean this case was moving well. Affordability and operational progress are different questions. We need both the risk flag and the evidence-based workflow."

If the initial request is already confirmed: show the recorded verifier and continue with the receipt step. **If receipt is already confirmed:** use a prepared recording of this same sequence or a rehearsed separate synthetic case. Never pretend that an already-completed transition just happened.

If AI takes too long: after roughly ten seconds say, "The manual evidence path is available while the model responds." Demonstrate Record An Update. Do not announce a live AI result when the source tag identifies a fallback.

10 — Before And After

Time: 16:15–18:00. **Screen:** page 10. Keep the case open in the adjacent tab.

Say:

"Let us compare exactly what changed. Before review, the confirmed case said 'with bank', while the more specific document issue existed in a message and an unconfirmed proposal. After review, the missing payslip became a recorded blocker with a responsible next action.

Before the receipt was verified, the system could not safely treat the document as received. After confirmation, the outstanding document cleared and the case returned to waiting for the bank decision. The confirming person and timestamps remain available.

We have demonstrated a change in operational state, not a measured reduction in processing time. We have not shown a recovered sale or a faster bank approval. Those are outcomes the pilot must measure.

The practical improvement is that the next person does not need to reconstruct the file from memory. They can see the blocker, the action, the evidence and what responsibility remains.

That is what the system causes: a specific follow-up and a verified handoff. The screen is the interface to that behaviour."

11 — Legal: An Approved Loan Still Needs A Signing

Time: 18:00–19:30. **Screen:** page 11, then [Legal](#) as **Arvind Raj · Legal Admin**. Use the top-right persona menu; the current menu includes the person's name.

Actions: locate **BK-9006 / A-23A-02 / Chen Jia Hao** under No Appointment Yet. Select **Record Appointment** to show the preset form. In a prepared synthetic rehearsal, record the appointment as arranged on 18 September with the appointment on 19 September, then show it under Appointment Set, Not Signed. During a shared-demo

walkthrough without a prepared state, show the form and cancel. Say explicitly which demonstration you performed.

Say:

"Our first case is still awaiting the bank. This second synthetic case has an approved loan, but no SPA appointment. It demonstrates the next handoff.

Legal's queue separates cases without an appointment from cases with an appointment that have not signed. The next action here is to arrange and record the appointment. The row opens the relevant update directly.

An appointment does not count as a signed SPA. Once Legal verifies execution, they record the signing, and the case leaves this waiting queue. The business metric counts that verified execution, not a scheduled date.

We also show panel workload, but the synthetic assignment of firms cannot establish which real solicitor performs better."

12 — Management: Cause An Internal Follow-Up

Time: 19:30–21:00. **Screen:** page 12, then [Manager](#) as **Project Manager · Manager**.

Actions: open **Suggestions**, then **Why This Case** on an overdue case. Show **Request Follow-Up From [name]**. In the prepared demo, select it and show **Follow-Ups Already Sent** or **Awaiting [name]**. Switch to that named recipient through the persona menu and show the task on `/chase`. Do not click multiple cases or duplicate an already-sent follow-up.

Say:

"The manager sees a cross-department view and the reason for the suggested intervention. Here the trigger is a wait beyond the configured expectation. These expectations are prototype settings, not measured service levels for Chin Hin's banks or solicitors.

The manager's button creates or flags an internal follow-up for the person responsible for the case. We can then open that person's Today page and see the work assigned to them. This is a concrete change in accountability.

The word 'sent' here means an internal task has been assigned. It does not mean an external banker received a WhatsApp message. Staff still use their existing approved communication channels.

Ask MortarAI is available for staff questions, but this presentation's proof is the case-and-task workflow. An answer is useful when it helps staff inspect the evidence and choose an appropriate action."

13 — Forecast: Explain The Estimate And Its Limits

Time: 21:00–23:00. **Screen:** page 13, then [Forecast](#) as **Project Manager**.

Actions: show expected signings, forecast range and bookings in scope. Expand **Stage Conversion Rates, How Well The Method Backtests** and **Assumptions**. Read the values actually visible; avoid memorising a demo

number. The inspected state showed approximately 21, range 17–25, 50 bookings.

Say:

“Management should not treat every unsigned booking as certain to convert. Mortar estimates the probability of a verified signing within thirty days of the original booking date. That is the outcome horizon used here; it is not a rolling forecast of cash arriving over the next thirty days.

The method looks at resolved cases that reached a similar stage at a similar booking age. It uses a broader comparison when a narrow group has too few cases. Then it adds the estimated probabilities across the live bookings. The range comes from repeated simulated outcomes under those estimates.

Stage rates show the sample size and uncertainty. The accuracy check looks back using the evidence available at the historical cutoff. It checks the method against what subsequently happened in the dataset.

There are two limitations. First, the history here is synthetic, so this check demonstrates the calculation rather than predictive performance on Chin Hin records. Second, the range reflects the model's assumptions. It does not capture every source of uncertainty or a market shock affecting many buyers at once.

During a pilot we would calibrate the rates on appropriate company cohorts, compare against a simple baseline and evaluate held-out outcomes. We would also reconcile signed SPAs with Legal.

Finance can use that as a more explicit conversion estimate. To estimate cash, we would additionally need payment schedules, billing milestones and actual collections. Multiplying expected SPAs by unit prices would not establish cash brought forward.”

14 — Technology Second

Time: 23:00–24:00. **Screen:** page 14. **Cue:** one record → shared rules → human verification → manual continuity.

Say:

“Now that the business workflow is visible, here is how we built it. The React interface presents the desks. A Bun server reads and writes the PostgreSQL case record. Shared TypeScript rules derive the stages, document blockers, responsibility and forecast, so the desks use consistent logic.

Jev handles structured interpretation and suggestions. Ask MortarAI is a separate assistant for staff questions. Confirmed evidence and statistical rules drive the operational state and forecast.

The server keeps model credentials out of the browser. A cached answer is labelled, and manual recording remains available. If an interpretation is wrong, staff can dispute it and record the correct evidence.

The important engineering choice is that the daily workflow can continue without a successful live model call.”

Do not claim: production SSO, a universally reliable three-second response, automatic WhatsApp ingestion or independent production security approval.

15 — Derive The Business Impact

Time: 24:00–26:00. **Screen:** page 15. **Cue:** explicit assumption → calculation → pilot replacement.

Say:

“We measure one primary outcome: the share of eligible cohort bookings that reach a verified signed SPA within thirty days of booking. Eligibility and the denominator must be defined before the pilot starts, so difficult cases are not quietly excluded after the result is known.

Here is an illustrative impact calculation, not a claim about Chin Hin. Suppose an audit identifies forty stalled bookings that could benefit from targeted follow-up. Suppose ordinary follow-up converts twenty percent within thirty days, and the assisted workflow converts forty percent.

Forty multiplied by the twenty-percentage-point difference gives eight additional verified SPAs. At a five-point difference, the same group gives two additional SPAs. At no difference, it gives none.

If those forty cases belong to a cohort of one hundred bookings, eight additional signings would increase the whole cohort's rate by eight percentage points. A twenty-point improvement in a subgroup is not a twenty-point improvement across the entire cohort.

Every input in this example is an assumption. We replace the eligible case count, baseline conversion and uplift with observed pilot data.

We also record how quickly unresolved units reach an authorised release decision and how much staff effort the process requires. These are supporting checks. The primary result remains verified SPA conversion, with an ordinary follow-up comparison and a clear observation window.

More dashboard visits would not establish success. More confirmed conversions would—and a credible comparison would tell us whether Mortar contributed.”

16 — How Adoption Actually Starts

Time: 26:00–28:00. **Screen:** page 16, then [Add Bookings](#) as **Project Manager** or **Nurul Aina · Sales Admin**. Select **Upload A Sheet** to show intake without submitting a file during the talk. Keep a rehearsed synthetic sheet ready if you want to demonstrate row review.

Say:

“Adoption begins with existing booking data. The working prototype supports typed intake and spreadsheet upload with a review step. For the pilot we would agree the export fields, booking identifiers and reconciliation process with the existing system owner. We would avoid a second competing master record.

In the first two weeks, one project and one Sales Administrator establish the baseline. They reconcile a sample of closed cases with Loan Administration and Legal, identify the main recoverable blocker and agree the first playbooks.

In weeks three and four, the administrator uses the daily queue in a short morning routine. They choose the next action, contact the relevant party through existing channels, then record and verify the response. The manager reviews overdue internal tasks. We measure whether this saves work or adds duplicate entry.

For the outcome comparison, we propose an intake cohort enrolled between days fifteen and forty-four, spanning the assisted-queue and live-operation periods. A comparable group continues ordinary follow-up. Where practical, assignment is randomised; otherwise, project, buyer mix, booking age and bank conditions need to be accounted for. We agree the comparison with the sponsor before starting.

Each booking then gets the full thirty-day observation period. That ends by day seventy-four, with Legal's reconciliation completed by day eighty-four. Weeks five to eight are the main intervention period; weeks nine to twelve complete the observation and audit.

Buyers, bankers and solicitors continue using their existing channels. The internal administrator owns the new routine. We seek integration only after the pilot shows where manual intake creates a real burden, and in alignment with Chin Hin's existing transformation programme."

17 — Built, Proposed And Still Unproven

Time: 28:00–29:15. **Screen:** page 17. **Cue:** demonstrated mechanism ≠ real-company impact.

Say:

"What exists today is the working case record, staff desks, task workflow, message proposals with human review, manual updates, legal queue, management follow-up and statistical forecast. The deployment uses synthetic records.

What remains proposed includes live-company calibration, production identity controls and approved integration with existing operational systems. Automated WhatsApp ingestion, document OCR and the more ambitious mortgage-rescue ideas are not part of the workflow we demonstrated.

The unknown is the business effect. We have not measured additional SPAs, inventory days saved or cash accelerated at Chin Hin. We have also not established the company's current internal leakage causes or system-level permissions.

Those boundaries determine the next step: a small, supervised pilot with a known owner, an agreed comparison and verified outcomes. Every team member should be able to explain that whole chain, including what the prototype does not yet prove."

18 — Close And Concrete Request

Time: 29:15–30:00. **Screen:** page 18. Stop operating the browser and look at the audience.

Say:

“The change we are asking you to evaluate is simple: a stalled booking gets a specific next action, a responsible person and a verified record of what happened. Management then sees a conversion estimate grounded in that record.

We have demonstrated the workflow on synthetic cases. We propose testing it on one suitable project with one administrator, approved data access and a twelve-week evaluation of thirty-day verified SPA conversion.

The three things we need to agree are the pilot owner, the source of booking and signing evidence, and the comparison against ordinary follow-up.

Thank you. We are ready to discuss the business logic, the workflow, the implementation and the parts that still need validation.”

Questions Everyone Should Be Able To Answer

Likely Question	Concise Answer	Page To Return To
Why does Chin Hin need this if it has ERP?	We have not established that its existing modules cannot do this. The pilot tests a specific follow-up gap and should fit its system of record. Kingdee is publicly announced; actual project modules must be confirmed.	3, 16
What exactly changes after a click?	A next action becomes an assigned task, or reviewed evidence changes the case's blocker and next responsibility. A manager's follow-up becomes internal work.	8, 9, 12
Can you solve a bank rejection?	We cannot change the bank decision. Staff can resolve missing documents, chase processing or review another application. Mortar supports that coordination.	4, 6
Is your survey representative?	No. Eight self-selected practitioners with mixed roles provide directional priorities, not Chin Hin's cause shares or conversion rates.	4
Have you interviewed Chin Hin staff?	The repository's practitioner interview template is blank. We do not present it as a completed interview. We have the brief, survey and public sources.	4
Have you increased conversion?	Not on company records. We demonstrated mechanics on synthetic data and proposed the comparison needed to measure impact.	15, 17
Where did the eight extra SPAs come from?	Forty assumed eligible cases multiplied by an assumed 20-percentage-point uplift. It is an illustration, not an observed result.	15
Does a cleared payslip mean a recovered sale?	No. It clears one verified blocker. The bank decision and signing remain separate milestones.	9, 10
Does Complete verify a milestone?	Completing a task records work completion. A milestone requires its own confirmed case event.	9
Is the forecast AI-generated?	No. It uses stage-and-age comparison rates and sums probabilities. The simulation and comparison assumptions still require company validation.	13
What does thirty days mean?	Within thirty days of the original booking date, both for the proposed success metric and the prototype's signing target. It is not a cash horizon.	13, 15
Does your range cover all uncertainty?	No. It is a conditional simulation range. Data quality, rate-estimation uncertainty and correlated market changes require additional validation.	13

What if the AI is wrong or unavailable?	Review or dispute the proposal, then record verified evidence manually. Cached/unavailable sources are labelled.	9, 14
Does the system automatically send WhatsApp?	No. Staff paste messages and use existing channels. Internal follow-up assignment is not external message delivery.	8, 12, 17
Is the 40% DSR a bank approval policy?	It is a configurable prototype advisory threshold, not a universal bank policy or credit decision. Rates and eligibility vary.	14, 17
Does one rejected application close the booking?	No. Applications are tracked separately. The overall case can still have an active or approved alternative.	6
Who cancels or releases a unit?	An authorised person under company policy. Mortar can record that decision; it does not acquire authority from an AI score.	6, 17
How do you avoid gaming the metric?	Define eligibility beforehand, keep difficult eligible cases in the denominator, give the whole cohort its observation window and reconcile executions with Legal.	15, 16
What if the pilot cohort is too small?	Report uncertainty and operational findings. Do not claim reliable causal improvement from a few cases; extend observation or aggregate comparable cohorts.	15, 16
Who operates and maintains it?	Propose a named internal process owner plus a named technical owner. Their staffing, support time and costs must be agreed; free prototype hosting is not a production cost estimate.	16, 18
Is it safe to load real buyer records now?	The demo is not a production rollout. Approved access, identity controls, processing arrangements and retention rules are required before live data enters it.	17

If you do not know: “We have not verified that yet. Our current assumption is [state it]. We would check [specific data/person] before relying on it.” Do not guess or pass every question to a specialist.

Good Questions To Ask Chin Hin

- Which system is the authoritative booking and unit record for the candidate pilot project, and how does it relate to the Kingdee transformation?
- Who owns a booking's follow-up at each handoff today, and who can approve cancellation or unit release?
- What evidence does Legal use to mark an SPA as executed, and can we reconcile that evidence with the booking identifier?
- Of the last closed cohort, which failure causes and recoverable delays are recorded reliably?
- Can we create a fair ordinary-follow-up comparison and give every enrolled booking its complete thirty-day observation window?
- Which part of this workflow would remove duplicate work for your administrator, and which part would add it?

These questions are part of understanding the business. Ask relevant ones during Q&A rather than claiming you already know the answers.

Rehearsal And Demo Preparation

- Open the new deck and the Render app in separate tabs. Preload the app to avoid a cold start. Use a readable zoom, silence notifications and hide unrelated tabs.

- Check `/api/health`; it should report a healthy database. A configured model flag does not prove every live model call will succeed.

- Verify BK-9001's task and evidence state. Rehearsals change the shared demo.

Prefer a dedicated disposable demo deployment for repeatable mutations. Do not delete shared demo data simply to make the script match.

- Have the buyer-reply text in a plain-text note. Confirm the proposer/source tag and the human verifier on screen.

- Prepare a short recording of the critical transition on synthetic data as a fallback. Say it is a recording if you use it.

- Rehearse named-profile switches: Nurul Aina, Arvind Raj and Project Manager.

Current access is scoped; different profiles can see different case counts.

- Rehearse the manual Documents Received path and the Legal appointment form.

- Know the demo's fixed 18 September reference date. Distinguish event date, recorded date and a future appointment date.

- Time each section with the cue sheet. Never pad by reading every row. If you finish early, use the deeper explanation prompts below.

- Every member performs the full critical loop once and explains the survey, forecast, impact calculation and pilot without notes.

- Meaningful contributions must be real: evidence collection, workflow design, implementation, validation or pilot design. Do not invent authorship. Someone joining can present the section matching their actual contribution, while keeping this universal script's timing and argument.

If A Section Finishes Early

Use these within its allocated time; do not extend the overall thirty minutes.

- **Company context:** explain unsigned booking versus contracted unbilled

sale, using one synthetic unit rather than another financial statistic.

- **Evidence:** explain why a small mixed-role survey supports priorities but cannot estimate company leakage shares.

- **Demo:** ask the audience to look at the verifier and source; show the full event history and one separate bank application.

- **Forecast:** explain why a sample of three comparable cases is fragile and why the engine backs off to a broader comparison.

- **Impact:** calculate the five-point uplift scenario aloud: $40 \times 0.05 = 2$.

- **Adoption:** walk through the administrator's actual morning sequence: queue, action, existing communication channel, receipt check, recorded event.

Mapping To The Existing Qualification Deck

The existing repository deck has 20 pages. Its page numbers are different. The new companion deck avoids presenting stale evidence and role labels.

New Page	Related Old Page	Important Difference
1	1	Business-led spoken opening
2	None	Explicit thirty-minute agenda
3	None	Official company context and Kingdee announcement
4	4, 10	Updated survey $n = 8$; no company cause-share claim
5	2, 3	Workflow before, without an unsupported measured baseline
6	5, 6, 7	Four actual roles; evidence-to-action logic
7	16	Fixed demo date; named-profile session caveat
8	8	Today, Quick View and existing-task branch
9	7, 9	Actual request/receipt confirmation loop
10	None	Explicit operational before/after
11	6	Legal Admin is Arvind Raj; Finance is a beneficiary
12	None	Current Manager follow-up workflow
13	11, 12	Correct thirty-days-from-booking target and synthetic limits
14	10, 15	Render deployment; separate Jev and Ask MortarAI

15	17	Uplift is illustrative; subgroup/cohort arithmetic explicit
16	18	Coherent day 15–44 enrollment and day 74/84 completion
17	19	Ask MortarAI exists; production identity/integrations still proposed
18	20	Working Render link and concrete pilot request

Old page 13's financing assumptions and page 14's playbooks are useful backup topics. Present the 40% DSR and other seed parameters as prototype assumptions, not universally binding bank rules. Avoid the old deck's blanket PDPA claim, outdated "no generative chatbot" line and unqualified legal exposure claims.